

# Natural Language Processing

Prof. Dr. Jannik Strötgen  
[jannik.stroetgen@h-ka.de](mailto:jannik.stroetgen@h-ka.de)

Summer 2024

**Hochschule Karlsruhe**  
University of  
Applied Sciences

Fakultät für  
**Informatik und**  
**Wirtschaftsinformatik**



# Preliminary Agenda

Date	Topic	Date	Topic
19.03.2024	Organisation & motivation	14.05.2024	Contextual Embeddings
26.03.2024	Introduction to NLP and	21.05.2024	<i>no lecture (Whitsun break)</i>
02.04.2024	<i>no lecture (Easter)</i>	28.05.2024	Text Mining and Sentiment Analysis
09.04.2024	Pre-Processing and Part-of-Speech Tagging	04.06.2024	Information Extraction & QA
16.04.2024	Parsing	11.06.2024	Applications exploiting NLP
23.04.2024	Named Entity Recognition and Disambiguation	18.06.2024	NLP with LLMs
30.04.2024	Similarity and Search	25.06.2024	My Research Topics
07.05.2024	Language Models: Static Word Embeddings	02.07.2024	Recap, exam preparation

# [3] Part-of-Speech Tagging

# Outline

## [3] Part-of-Speech Tagging

1. Parts of speech (POS)
2. Tagging
3. Feature-based POS tagging
4. Probabilistic POS tagging
  - Unigram tagger
  - N-gram tagger
5. Rule-based POS tagging
  - Brill tagger

# From the Introduction

## 2. Part-of-Speech (POS) Tagging

Adding information about word types helps in determining the roles that words play and in understanding the meaning of a text. For example, we can identify verbs, adverbs, nouns, determiner (articles), etc.

Ah, distinctly I remember it was in the bleak December;  
And each separate dying ember wrought its ghost upon the floor.  
Eagerly I wished the morrow;—vainly I had sought to borrow  
From my books surcease of sorrow—sorrow for the lost Lenore—  
For the rare and radiant maiden whom the angels name Lenore—  
Nameless here for evermore.



# Parts of Speech (POS)

# Let's Disambiguate Some Homonyms!

The **die** is **fair**  
but the house wins.



You will **die** at the **fair**  
if you enter the house.

# Homonyms and Parts of Speech



The **die** is **fair** but the house wins.  
noun      adjective



You will **die** at the **fair** if you enter...  
verb      noun

# Homonyms and Parts of Speech

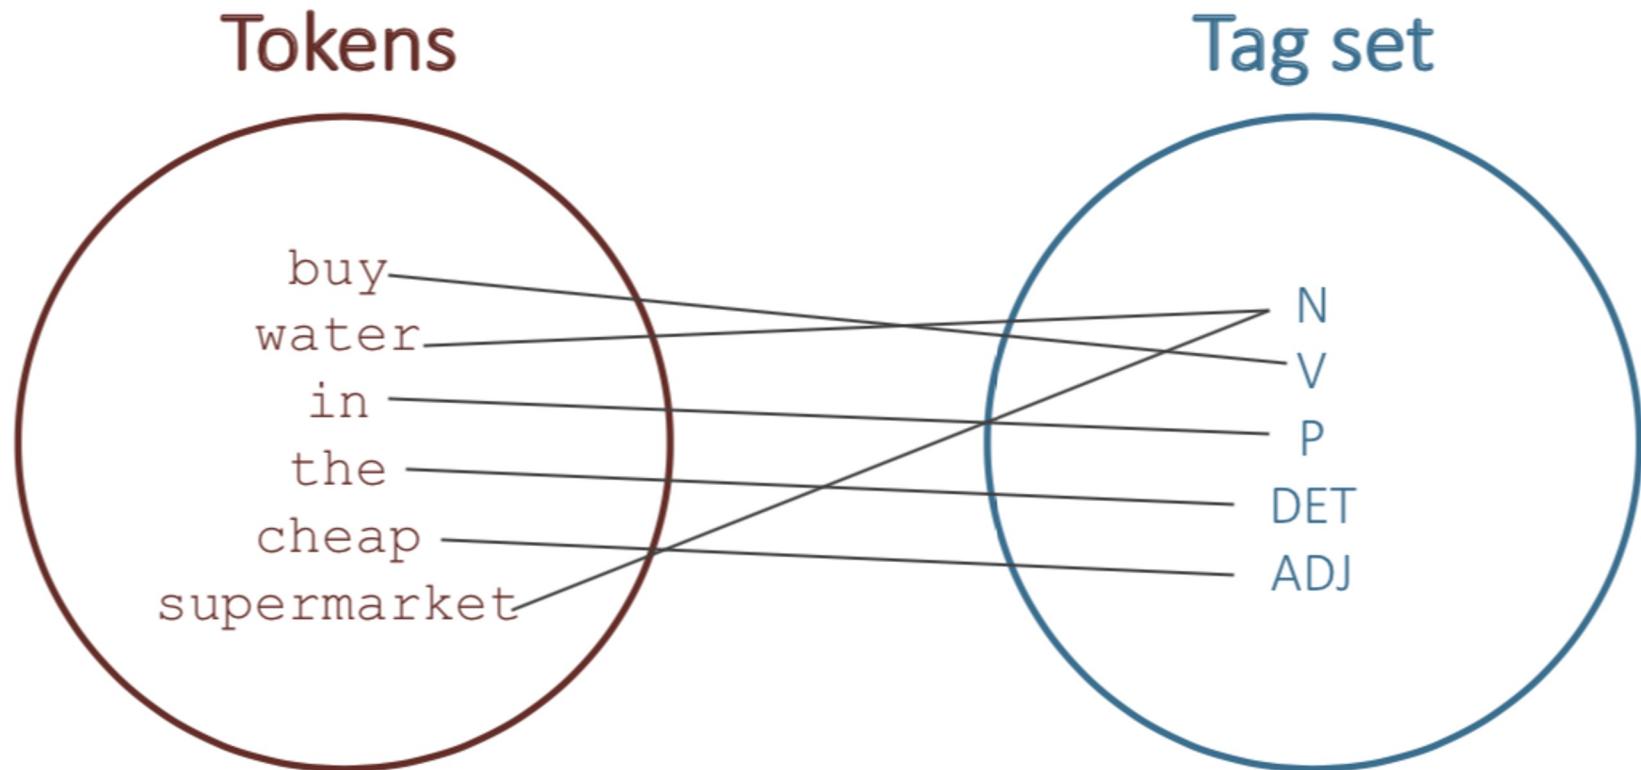


The **die** is **fair** but the house wins.  
DT NN VBZ JJ CC DT NN VBZ

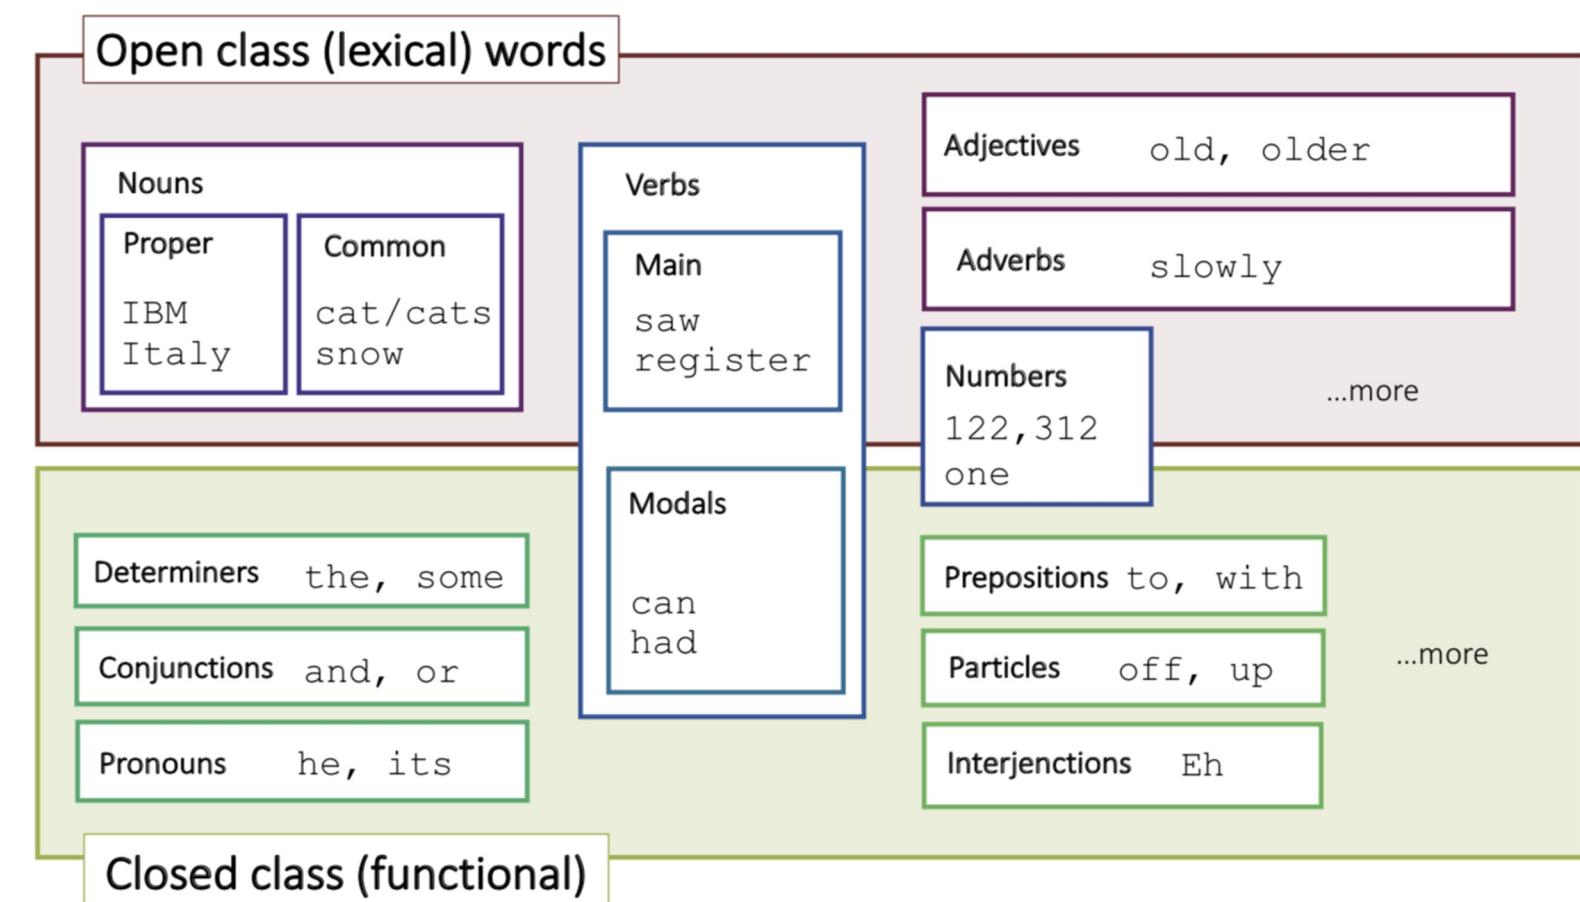


You will **die** at the **fair** if you enter...  
PRP MD VB IN DT NN IN PRP VBP

# Part of Speech Tagging



# Part of Speech Hierarchy



# Part of Speech Tagging Task

Words often have more than one POS

- The **die** is cast → **NN**
- To **die** or not to die → **VB**
- The POS tagging problem addresses how to determine the POS tag for a particular instance of a word.

Example:

- Input:      The                  die                  is                  fair
- Ambiguity: DT                  NN/VB              VBZ              JJ/NN
- Output:     The [DT] die [NN] is [VBZ] fair [NN]

# Tagging

# The Purpose of POS Tagging

1. Collapsing distinctions into equivalence classes
  - All personal pronouns (You, they, etc.) tagged with PRP
2. Introducing distinctions to resolve ambiguities
  - E.g., leaves could be tagged with NN or VB
3. Helpful features in classification and prediction tasks
  - **Named entity extraction:**  
Essen could be the name of a German city, but not if it is tagged VB
  - **Machine translation:**  
Are translated sentences syntactically correct?  
(stay tuned for syntactic parsing next week!)

# Terminology

## Tagging

- The process of associating labels with each token in a text
- Not limited to parts of speech (e.g., helpful in sentiment analysis)

## Tags

- The labels that are assigned to the tokens

## Tag Set

- The collection of tags that is used for a particular task
  - Tag sets are defined by linguists and not unambiguous.
  - Different languages require different language-specific tag sets
- It is difficult to compare tags between different languages

# Tagset Requirements

**Informativeness**-related parameters:

- The size of the tagset
  - The degree of the ambiguity
- Does the tagset represent all unique cases in the language, or does ambiguity still exist?

**Specifiability**

- Inter-annotator (interjudge) agreement should be good
- What is the agreement when multiple (trained) linguists tag the same text?

# Tagset Requirements: Example

Collapsing parts of the Penn Treebank tagset:

- JJ: adjective example: smart
- JJR: adjective (comparative) example: smarter
- JJS: adjective (superlative) example: smartest
- Collapse all three tags into one: JJ

What happens to...

- informativeness? Size of the tagset decreases. Ambiguity increases.
- specifiability? Inter-annotator agreement likely increases.

# Frequently Used Tagset

English:

- Penn Treebank: 45 tags
- Brown corpus: 87 tags
- Lancaster UCREL C5: 61 tags (used to tag the BNC)
- Lancaster C7: 145 tags

German:

- Stuttgart-Tübingen Tagset (STTS): 54 tags

# Penn Treebank Tagset

Number	Tag	Description
1.	CC	Coordinating conjunction
2.	CD	Cardinal number
3.	DT	Determiner
4.	EX	Existential there
5.	FW	Foreign word
6.	IN	Preposition or subordinating conjunction
7.	JJ	Adjective
8.	JJR	Adjective, comparative
9.	JJS	Adjective, superlative
10.	LS	List item marker
11.	MD	Modal
12.	NN	Noun, singular or mass
13.	NNS	Noun, plural
14.	NNP	Proper noun, singular
15.	NNPS	Proper noun, plural
16.	PDT	Predeterminer
17.	POS	Possessive ending
18.	PRP	Personal pronoun

Number	Tag	Description
19.	PRP\$	Possessive pronoun
20.	RB	Adverb
21.	RBR	Adverb, comparative
22.	RBS	Adverb, superlative
23.	RP	Particle
24.	SYM	Symbol
25.	TO	to
26.	UH	Interjection
27.	VB	Verb, base form
28.	VBD	Verb, past tense
29.	VBG	Verb, gerund or present participle
30.	VBN	Verb, past participle
31.	VBP	Verb, non-3rd person singular present
32.	VBZ	Verb, 3rd person singular present
33.	WDT	Wh-determiner
34.	WP	Wh-pronoun
35.	WP\$	Possessive wh-pronoun
36.	WRB	Wh-adverb

# Stuttgart-Tübingen Tagset (STTS)

ADJA	attributives Adjektiv	[das] grosse [Haus]	PPER	irreflexives Personalpronomen	ich, er, ihm, mich, dir
ADJD	adverbiales oder prädikatives Adjektiv	[er fährt] schnell [er ist] schnell	PPOSSE	substituierendes Possessivpronomen	meins, deiner
ADV	Adverb	schon, bald, doch	PPOSAT	attribuierendes Possessivpronomen	mein [Buch], deine [Mutter]
APPR	Präposition; Zirkumposition links	in [der Stadt], ohne [mich]	PRELS	substituierendes Relativpronomen	[der Hund ,] der
APPRART	Präposition mit Artikel	im [Haus], zur [Sache]	PRELAT	attribuierendes Relativpronomen	[der Mann ,] dessen [Hund]
APPO	Postposition	[ihm] zufolge, [der Sache] wegen	PRF	reflexives Personalpronomen	sich, einander, dich, mir
APZR	Zirkumposition rechts	[von jetzt] an	PWS	substituierendes Interrogativpronomen	wer, was
ART	bestimmter oder unbestimmter Artikel	der, die, das,	PWAT	attribuierendes	welche [Farbe],
CARD	Kardinalzahl	ein, eine, ...	PWAV	Interrogativpronomen	wessen [Hut]
FM	Fremdsprachliches Material	zwei [Männer], [im Jahre] 1994	PAV	adverbiales Interrogativ- oder Relativpronomen	warum, wo, wann,
ITJ	Interjektion	[Er hat das mit "]	PTKZU	Pronominaladverb	worüber, wobei
ORD	Ordinalzahl	A big fish [ " übersetzt]	PTKNEG	"zu" vor Infinitiv	dafür, dabei, deswegen, trotzdem
KOUI	unterordnende Konjunktion mit "zu" und Infinitiv	mhm, ach, tja	PTKVZ	Negationspartikel	zu [gehen]
KOUS	unterordnende Konjunktion mit Satz	[der] neunte [August]	PTKANT	abgetrennter Verbzusatz	nicht
KON	nebenordnende Konjunktion	um [zu leben],	PTKA	Antwortpartikel	[er kommt] an, [er fährt] rad
KOKOM	Vergleichskonjunktion	anstatt [zu fragen]	SGML	Partikel bei Adjektiv	ja, nein, danke, bitte
NN	normales Nomen	weil, dass, damit,	SPELL	oder Adverb	am [schönsten],
NE	Eigennamen	wenn, ob	TRUNC	SGML Markup	zu [schnell]
PDS	substituierendes Demonstrativpronomen	und, oder, aber	VVFIN	Buchstabierfolge	S-C-H-W-E-I-K-L
PDAT	attribuierendes Demonstrativpronomen	als, wie	VVIMP	Komposition-Erstglied	An- [und Abreise]
PIS	substituierendes Indefinitpronomen	Tisch, Herr, [das] Reisen	VVINF	finites Verb, voll	[du] gehst, [wir] kommen [an]
PIAT	attribuierendes Indefinit-pronomene ohne Determiner	Hans, Hamburg, HSV	VVIZU	Imperativ, voll	komm [!]
PIDAT	attribuierendes Indefinit-pronomene mit Determiner	dieser, jener	VVPP	Infinitiv mit "zu", voll	gehen, ankommen
		jener [Mensch]	VAFIN	Partizip Perfekt, voll	anzukommen, loszulassen
		keiner, viele, man, niemand	VAIMP	finites Verb, aux	gegangen, angekommen
		kein [Mensch],	VAINF	Imperativ, aux	[du] bist, [wir] werden
		irgendein [Glas]	VAPP	Infinitiv, aux	sei [ruhig !]
		[ein] wenig [Wasser],	VMFIN	Partizip Perfekt, aux	werden, sein
		[die] beiden [Brüder]	VMINF	finites Verb, modal	gewesen
			VMPP	Infinitiv, modal	dürfen
			XY	Partizip Perfekt, modal	wollen
				Nichtwort, Sonderzeichen	gekonnt, [er hat gehen] können
				enthaltend	3:7, H2O,
				\\$,	D2WX3
				\\$.:	,
				\\$(	. ? ! ; :
					- [ , ]

# Performance: How Difficult is POS Tagging?

How many tags can algorithms correctly predict (tag accuracy)?

- With current methods about 97.85% ([https://aclweb.org/aclwiki/POS\\_Tagging\\_\(State\\_of\\_the\\_art\)](https://aclweb.org/aclwiki/POS_Tagging_(State_of_the_art)))

But a simple naïve baseline can reach 90% already:

- Tag every token with its most frequent tag
- Tag unknown tokens as nouns

This simple method performs well because:

- Many words are unambiguous
- The algorithm “gets points” for unambiguous tokens, e.g.
- Articles such as the, a, etc.
- Punctuation marks

# Feature-based Tagging

# Word-level Features for Tag Prediction

We can obtain a lot of information from a word itself to create tagging rules:

- Word              the              the → DT
- Prefixes:          unhealthy        un- → JJ
- Suffixes:          importantly     -ly → RB
- Capitalization:   Meridian        CAP → NNP

# Tag Sequence Features for Tag Prediction

Knowledge of tag combinations for neighboring words enables us to leverage syntactic rules to make decisions.

Example:

- Input: Bill saw that man yesterday
- Sequence 1: NP NN DT NN NN
- Sequence 2: VB VBD IN VB NN

Syntactically impossible sequence

# POS Performance Improvement (and Why It Matters)

Using only word-level and tag sequence features:

- We can reach 93.7% tag accuracy already
- “Woah, we’re **half way** there”  
(on the way to state-of-the-art performance)

So why should we care about the remaining 6.3%?

- NLP happens in what we call “The NLP Pipeline”
- Error propagation is a critical issue



# Towards More Advanced POS Tagging

## Probabilistic methods

- Use a **corpus for training**
- Obtain POS probabilities from **manual annotations** of the corpus
- Language-agnostic: Method can be applied to different languages

## Rule-based methods

- Use linguistic knowledge to solve ambiguous cases
- Knowledge is mapped to rules
- Rules are easy to interpret
- Each language needs its own rules

# Probabilistic Tagging

# Tagging With Lexical Frequencies

**Problem:** How can we determine the most likely POS tag for a given token?

**Solution:** Derive probabilities from a large, annotated text corpus

- Count POS tag frequencies for a given token in the annotated corpus
- Select the most probable (= frequent) POS tag in new, unseen data

# Tagging With Lexical Frequencies

Example:

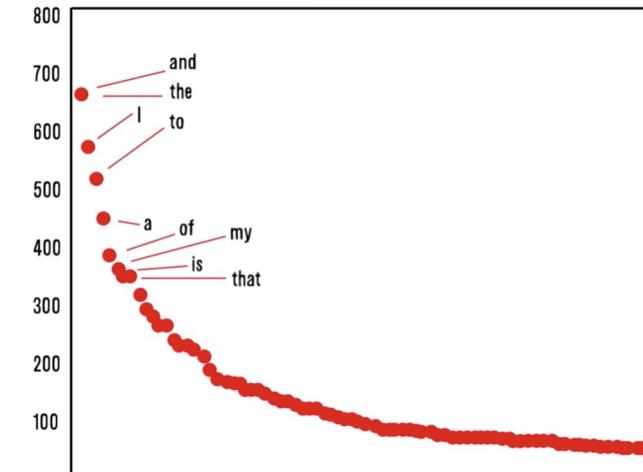
- Sebastian [NNP] Vettel [NNP] is [VBZ] expected [VBN] to [TO] **race** [VB] tomorrow [NN].
- People [NNS] continue [VBP] to [TO] inquire [VB] the [DT] reason [NN] for [IN] the [DT] **race** [NN] to [IN] outer [JJ] space [NN].
- To assign a tag to `race`, given its lexical frequency, we choose the tag that has the higher probability:  $p(\text{VB}|\text{race})$  or  $p(\text{NN}|\text{race})$ ?

# Unigram Tagger: A Naïve Prediction Model

- How a unigram tagger works:
- We have access to an **annotated text corpus** (= our training data)
- We create **statistics** of how many times each token is seen with each POS tag (= the training process)
- Based on these learned frequencies, we use the trained model on new, **unseen text** to associate with each word its most likely POS tag (= prediction)

Is it really this easy?

- Unfortunately, no!
- **Problem:** many words might never be seen in the training data (out-of-vocabulary)
- **Solution:** specify a default tag as “backoff”



# Unigram Tagger Example: Training

Training data:

Consider that result that George has achieved. That was simply amazing, although, he did not learn anything.

Manual tagging of the training data produces:

Consider[VB] that[DT] result>NN] that[PRP] George>NNP]  
has[VB] achieved[VB] .[SYM] That[PRP] was[VB] simply[RB]  
amazing[RB] , [SYM] although[CC] , [SYM] he[PRP] did[VB]  
not[RB] learn[VB] anything>NN] .[SYM]

(For simplicity, all verb forms are collapsed to VB in the example)

# Unigram Tagger Example: Prediction

Input test sentence:

George achieved that result.

Tagging of input according to the training data:

George [NNP] achieved [VB] **that [k]** result [NN] . [SYM]

Prediction:

- George, achieved, result, and . are unambiguous and can be assigned a POS tag directly
- k = [ DT ] | [ PRP ] requires a decision based on the training data

# Unigram Tagger Example: Prediction Probabilities

Manually tagged training data:

Consider [VB] that [DT] result [NN] that [PRP] George [NNP]  
has [VB] achieved [VB] . [SYM] That [PRP] was [VB] simply [RB]  
amazing [RB] , [SYM] although [CC] , [SYM] he [PRP] did [VB]  
not [RB] learn [VB] anything [NN] . [SYM]

Test input:

George [NNP] achieved [VB] that [k] result [NN] . [SYM]

POS probabilities according to the model trained on the training data:

$$p(k=DT | \text{that}) = 1/3$$

$$p(k=PRP | \text{that}) = 2/3$$

→ the model predicts that[PRP]

Unfortunately,  
[ PRP ] is wrong!

# Unigram Tagger: The Fundamental Problems

- Why is the unigram tagger not working?
- We are not computing probabilities properly: use Bayes' rule instead!  
Among all possible tags  $k'$ , the best choice for  $k=k'$  maximizes:

$$p(k = k' | \text{that}) = \frac{p(\text{that} | k = k') * p(k')}{p(\text{that})}$$

constant (can be ignored)

But, more simply: The most frequent tag is obviously not always correct!

- We need to know in which sense **that** is being used.
- We need to take the **context** into account.
- George [NNP] achieved [VB] **that[k]** **result[NN]** . [SYM]

# N-grams: Encoding Cooccurrence Context

An **n-gram** is a contiguous sequence of n items from a given sample of text.

- The n stands for how many terms are used
- n = 1: **unigram**, one token (0th order)
- n = 2: **bigram**, two tokens (1st order)
- n = 3: **trigrams**, three tokens (2nd order)
- etc., but n > 5 is rare (take a minute to figure out why)

Example:

- Input: George [NNP] achieved [VB] that [DT] result [NN]
- Token bigrams: George achieved; achieved that; that result
- Tag bigrams: NNP VB; VB DT; DT NN

# Using N-grams

N-grams can encode different types of contexts, e.g.,

- Character based n-grams
- Token-based n-grams
- POS-based n-grams
- Combinations of the above

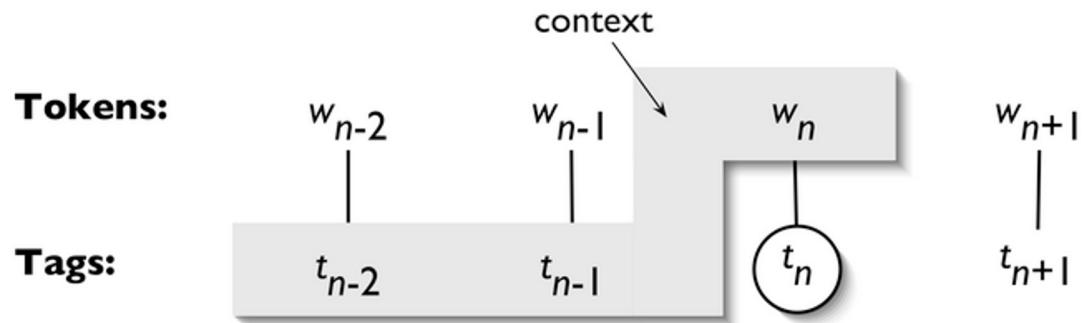
Ordering and context in n-grams:

- We often define context to mean “adjacent”, but this is not required
- We use n-grams to help us determine the context in which some linguistic phenomenon happens, e.g.,
  - Words before and after a period to determine if it is the end of a sentence
  - Detecting **open compound nouns** (e.g., outer space)

# N-gram POS tagger

Core idea:

- Use the unigram probability for the current token ( $\rightarrow$  unigram tagger)
- Use the predicted preceding  $n-1$  tags
- Optionally: We could also use the preceding  $n-1$  tokens, but this is impractical



Bigram tagger:

- Prediction: What is the most likely tag for token  $n$ , given token  $n$  and tag  $n-1$ ?
- The tagger picks the tag which is most likely, given that context

# Bigram Tagger Example

Back to our manually tagged training data:

Consider [VB] **that**[DT] result[NN] **that**[PRP] George[NNP]  
has[VB] achieved[VB] .[SYM] **That**[PRP] was[VB] simply[RB]  
amazing[RB] , [SYM] although[CC] , [SYM] he[PRP] did[VB]  
not[RB] learn[VB] anything[NN] .[SYM]

Test input:

George[NNP] achieved[VB] that[k] result[NN] .[SYM]

Preceding the token **that** is a token with the tag [VB]. Thus, we obtain:

$$p(k=DT|that) = p(\mathbf{k}=DT|VB) * p(that|DT) = 1/6 * 1/1 = 1/6$$

$$p(k=PRP|that) = p(\mathbf{k}=PRP|VB) * p(that|PRP) = 0/6 * 2/3 = 0$$

→ the model selects **that** [DT] as prediction



# Combining Multiple Taggers

To improve tagging accuracy, we should use more accurate algorithms when we can and simpler algorithms with broader coverage when needed.

- Attempt to tag the token with the 1st order tagger (e.g., bigram tagger)
- If the 1st order tagger is unable to find a tag for the token, fall back to using a 0th order tagger (e.g., unigram tagger)
- If the 0th order tagger is also unable to find a tag, use the default tagger to find a tag (e.g., feature based tagger)

Important note:

Bigram and trigram taggers need the previous tag context to assign new tags. If they see a [na] tag in the previous context, they will also output [na].

→ Error propagation

# Rule-based Tagging

# Rule-based Tagging

Valid reasons for linguistic complaints? Yes!

- Languages are more than cooccurrences statistics!
- This is just a heap of numbers and probabilities!
- Where is the linguistic knowledge?

Solution:

- We can use **handcrafted sets of rules** to tag input sentences
- For example, if a token follows a determiner, tag it as a noun

# The Brill Tagger

An example of Transformation-based Learning

- Core idea:
  - Use a probabilistic approach first
  - Then revise and correct it using contextual rules
- Similar to painting: sketch first, then paint properly
- Very popular (freely available and works fairly well)
- The rules are linguistically interpretable
- It is a supervised method and therefore requires a tagged corpus
- Example: [http://cst.dk/online/pos\\_tagger/uk/index.html](http://cst.dk/online/pos_tagger/uk/index.html)

# Brill Tagger: Process

## Initialization:

- Tag each word with the most likely POS (as observed in the corpus)
- For words that are not in the corpus:
  - If it is capitalized: label as noun
  - For tokens that end in the same three letters as a word with known tag: use the tag of the known word (e.g., label **ing** [VBG])

## Patching:

- Find suitable instances and apply patches to improve the result
- Example of a patch:
  - If one of the two preceding words is tagged as a determiner, change the tag of tokens that are tagged as a verb to noun

# Brill Tagger: Example

Input: Ask congress to increase grants.

Gold standard: VB NP TO VB NNS

Unigram tagger: VB NP TO **NN** NNS

Patch: Replace [NN] with [VB] when the previous tag is [TO]

# Brill Tagger: Patch Templates

- The preceding (following) token is tagged  $z$
- The token two before (after) is tagged  $z$
- One of the two preceding (following) tokens is tagged  $z$
- One of the three preceding (following) tokens is tagged  $z$
- The preceding (following) word is tagged  $z$  and the following token is tagged  $w$
- The preceding (following) word is tagged  $z$  and the word two before (after) is tagged  $w$
- The current word is (is not) capitalized
- The previous word is (is not) capitalized

# Brill Tagger: Patches

To generate patching rules:

- Compare initial tags with correct tags, setup a list of tagging error triples: (correct tag, incorrect tag, frequency)
  - For each error triple, find the patch template whose application results in the greatest error reduction.
  - For each error triple, compute the reduction in error that results from applying the patch (subtracting the number of new errors...).
  - Add the patch that results in the greatest improvement to the list of patches and apply it to the corpus.
  - Repeat until no further improvement can be gained.
- Result: tagging procedure (ordered list of transformations), which can be applied to new, untagged text.

# POS Tagging – Today (Example)

- Neural network-based methods are common for pos tagging with sota performance
- Either Bi-LSTMs or transformers can be used

10.12305v2 [cs.CL] 29 Oct 2021

**FAME: Feature-Based Adversarial Meta-Embeddings for Robust Input Representations**

Lukas Lange<sup>1,2,3</sup>   Heike Adel<sup>1</sup>   Jannik Strötgen<sup>1</sup>   Dietrich Klakow<sup>2</sup>

<sup>1</sup> Bosch Center for Artificial Intelligence, Renningen, Germany  
<sup>2</sup> Spoken Language Systems (LSV), Saarland University, Saarbrücken, Germany  
<sup>3</sup> Saarbrücken Graduate School of Computer Science, Saarbrücken, Germany  
{lukas.lange, heike.adel, jannik.stroetgen}@de.bosch.com  
dietrich.klakow@lsv.uni-saarland.de

**Abstract**

Combining several embeddings typically improves performance in downstream tasks as different embeddings encode different information. It has been shown that even models using embeddings from transformers still benefit from the inclusion of standard word embeddings. However, the combination of embeddings of different types and dimensions is challenging. As an alternative to attention-based meta-embeddings, we propose feature-based adversarial meta-embeddings (FAME) with an attention function that is guided by features reflecting word-specific properties, such as shape and frequency, and show that this is beneficial to handle subword-based embeddings. In addition, FAME uses adversarial training to optimize the mappings of differently-sized embeddings to the same space. We demonstrate that FAME works effectively across languages and domains for sequence labeling and sentence classification, in particular in low-resource settings. FAME sets the new state of the art for POS tagging in 27 languages, various NER settings and question classification in different domains.

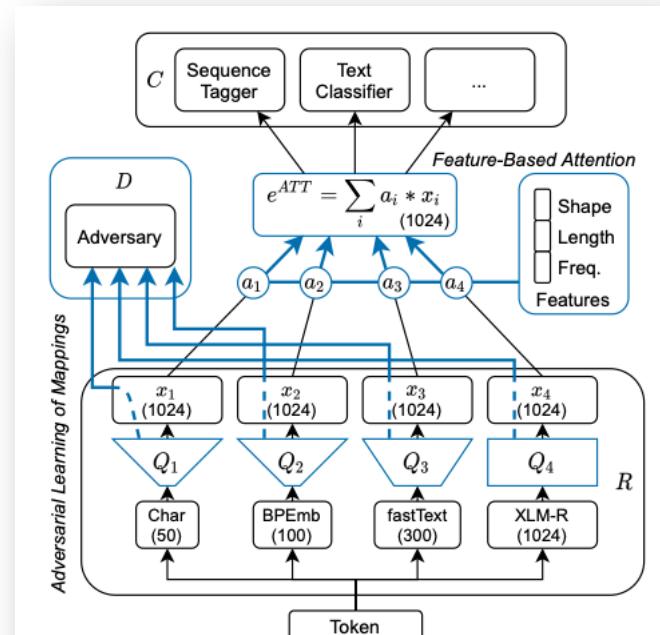


Figure 1: Overview of the FAME model architecture. Blue lines highlight our contributions.  $C$  (classifier),  $D$  (discriminator) and  $R$  (input representation) denote the components of adversarial training. The dimensions of intermediate representations are given in parentheses.

# POS Tagging – Today (Example)

- POS tagging results on high and low resource languages are quite good nowadays.
- In our work, good input representation was the key!
- We will learn about “input representation” and “embeddings” in later lectures.

	SOTA1	SOTA2	SOTA3	FAME
Bg (Bulgarian)	97.97	98.53	98.7	<b>99.53</b>
Cs (Czech)	98.24	98.81	98.9	<b>99.33</b>
Da (Danish)	96.35	96.74	97.0	<b>99.13</b>
De (German)	93.38	94.35	94.0	<b>95.95</b>
En (English)	95.17	95.82	95.6	<b>98.09</b>
Es (Spanish)	95.74	96.44	96.5	<b>97.75</b>
Eu (Basque)	95.51	94.71	95.6	<b>97.66</b>
Fa (Persian)	97.49	97.51	97.1	<b>98.68</b>
Fi (Finnish)	95.85	95.40	94.6	<b>98.67</b>
Fr (French)	96.11	96.63	96.2	<b>97.19</b>
He (Hebrew)	96.96	97.43	96.6	<b>98.00</b>
Hi (Hindi)	97.10	97.21	97.0	<b>98.35</b>
Hr (Croatian)	96.82	96.32	96.8	<b>97.96</b>
Id (Indonesian)	93.41	94.03	93.4	<b>94.24</b>
It (Italian)	97.95	98.08	98.1	<b>98.82</b>
Nl (Dutch)	93.30	93.09	93.8	<b>94.74</b>
No (Norwegian)	98.03	98.08	98.1	<b>99.16</b>
Pl (Polish)	97.62	97.57	97.5	<b>99.05</b>
Pt (Portuguese)	97.90	98.07	98.2	<b>98.86</b>
Sl (Slovenian)	96.84	98.11	98.0	<b>99.44</b>
Sv (Swedish)	96.69	96.70	97.3	<b>99.17</b>
Avg.	96.40	96.65	96.6	<b>98.08</b>
El (Greek)	-	98.24	97.9	<b>98.89</b>
Et (Estonian)	-	91.32	92.8	<b>97.07</b>
Ga (Irish)	-	91.11	91.0	<b>94.27</b>
Hu (Hungarian)	-	94.02	94.0	<b>97.72</b>
Ro (Romanian)	-	91.46	89.7	<b>96.64</b>
Ta (Tamil)	-	83.16	88.7	<b>91.10</b>
Avg.	-	91.55	92.4	<b>95.95</b>

Table 4: POS tagging results (accuracy) (using gold segmentation). SOTA1 refers to results from Plank et al. (2016), SOTA2 to Yasunaga et al. (2018) and SOTA3 to Heinzerling and Strube (2019). As Yasunaga et al. (2018), we split into high-resource (top) and low-resource languages (bottom).

# Further-Watching Material

Introduction to N-grams

<https://youtu.be/O7k8M8FwGLg?list=PLoROMvodv4rOFZnDyrIW3-nI7tMLtmijZ>

Estimating N-gram Probabilities

<https://youtu.be/JfOdK0mYtql?list=PLoROMvodv4rOFZnDyrIW3-nI7tMLtmijZ>

Thank you for your attention!

Questions?